


HireIQ

Multimodal AI Interview Evaluation Report

Candidate	Meera Iyer	Interview	AI Engineer Screening — Demo
Email	meera.demo@hireiq.local	Rubric	technical
Overall	68.0/100	Recommendation	Hold

HireIQ Skill Radar



Score Interpretation

Scores are evidence-based outputs of the configured rubric and LLM evaluation. Confidence and sentiment are supplementary signals and should not be used as proxies for protected traits.

Per-question Evaluation

Q1. Explain the difference between a process and a thread, and when you would prefer one over the other.

Technical	69	Problem Solving	71
Communication	68	Confidence	71
Overall	69	Recommendation	Hold

Transcript

A process has its own address space, while threads share memory inside a process. I would choose threads for lightweight concurrent work and processes for stronger isolation.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=Explain, the, difference, between, a

Q2. How would you design a REST API for a simple task management application?

Technical	66	Problem Solving	68
Communication	65	Confidence	68
Overall	66	Recommendation	Hold

Transcript

I would define resources and HTTP methods first, add request validation and authentication, and persist tasks in a relational database with pagination and clear error responses.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=How, would, you, design, a

Q3. You have an $O(n^2)$ algorithm over one million records. How would you investigate and improve it?

Technical	72	Problem Solving	74
Communication	71	Confidence	74
Overall	72	Recommendation	Hold

Transcript

I would profile the bottleneck first, confirm the complexity, and then look for a better algorithm or indexing strategy. For one million records I would avoid repeated nested scans and verify the improvement with measurements.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=You, have, an, $O(n^2)$, algorithm

Q4. What is overfitting in machine learning, and give two practical ways to reduce it?

Technical	63	Problem Solving	65
Communication	62	Confidence	65
Overall	63	Recommendation	Hold

Transcript

Overfitting happens when a model captures training noise and performs poorly on unseen data. I would use regularization and validation, and I could also add more data or simplify the model.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=What, is, overfitting, in, machine

Q5. Tell us about a technical problem you solved and how you communicated the trade-offs to your team.

Technical	70	Problem Solving	72
Communication	69	Confidence	72
Overall	70	Recommendation	Hold

Transcript

I had a slow data pipeline, profiled it, changed a repeated lookup into a cached structure, and explained the throughput versus memory trade-off to the team with measurements.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=Tell, us, about, a, technical

AI Ethics & Bias Note

HireIQ is a decision-support system, not an autonomous hiring authority. Reviewers should validate evidence, monitor for disparate effects, and avoid treating accent, sentiment, or confidence heuristics as measures of candidate worth.