


HireIQ

Multimodal AI Interview Evaluation Report

Candidate	Priya Rao	Interview	AI Engineer Screening — Demo
Email	priya.demo@hireiq.local	Rubric	technical
Overall	83.0/100	Recommendation	Proceed

HireIQ Skill Radar



Score Interpretation

Scores are evidence-based outputs of the configured rubric and LLM evaluation. Confidence and sentiment are supplementary signals and should not be used as proxies for protected traits.

Per-question Evaluation

Q1. Explain the difference between a process and a thread, and when you would prefer one over the other.

Technical	84	Problem Solving	86
Communication	83	Confidence	86
Overall	84	Recommendation	Proceed

Transcript

A process has its own address space, while threads share memory inside a process. I would choose threads for lightweight concurrent work and processes for stronger isolation.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=Explain, the, difference, between, a

Q2. How would you design a REST API for a simple task management application?

Technical	81	Problem Solving	83
Communication	80	Confidence	83
Overall	81	Recommendation	Proceed

Transcript

I would define resources and HTTP methods first, add request validation and authentication, and persist tasks in a relational database with pagination and clear error responses.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=How, would, you, design, a

Q3. You have an $O(n^2)$ algorithm over one million records. How would you investigate and improve it?

Technical	87	Problem Solving	89
Communication	86	Confidence	89
Overall	87	Recommendation	Proceed

Transcript

I would profile the bottleneck first, confirm the complexity, and then look for a better algorithm or indexing strategy. For one million records I would avoid repeated nested scans and verify the improvement with measurements.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=You, have, an, $O(n^2)$, algorithm

Q4. What is overfitting in machine learning, and give two practical ways to reduce it?

Technical	78	Problem Solving	80
Communication	77	Confidence	80
Overall	78	Recommendation	Proceed

Transcript

Overfitting happens when a model captures training noise and performs poorly on unseen data. I would use regularization and validation, and I could also add more data or simplify the model.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=What, is, overfitting, in, machine

Q5. Tell us about a technical problem you solved and how you communicated the trade-offs to your team.

Technical	85	Problem Solving	87
Communication	84	Confidence	87
Overall	85	Recommendation	Proceed

Transcript

I had a slow data pipeline, profiled it, changed a repeated lookup into a cached structure, and explained the throughput versus memory trade-off to the team with measurements.

Strengths: Correct core concept; Explains a practical trade-off; Evidence is relevant

Improvements: Add one more implementation detail; Quantify the example where possible

Rationale: Synthetic demo evaluation aligned to the configured rubric.

NLP: sentiment=positive; keywords=Tell, us, about, a, technical

AI Ethics & Bias Note

HireIQ is a decision-support system, not an autonomous hiring authority. Reviewers should validate evidence, monitor for disparate effects, and avoid treating accent, sentiment, or confidence heuristics as measures of candidate worth.